

time-series-gru-bike-and-rf

June 3, 2025

```
[1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from tensorflow.keras.layers import GRU, Bidirectional, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.regularizers import l2

# Load your dataset
train_df = pd.read_csv("/kaggle/input/bike-sharing-demand/train.csv")
test_df_raw = pd.read_csv("/kaggle/input/bike-sharing-demand/test.csv")
```

A Gated Recurrent Unit (GRU) is a type of Recurrent Neural Network (RNN) designed to capture dependencies in sequential data. It introduces gating mechanisms—specifically, the update and reset gates—that regulate the flow of information, allowing the model to retain or discard information as needed. This architecture addresses challenges like the vanishing gradient problem, enabling GRUs to model long-term dependencies more effectively than traditional RNNs.

Key Characteristics of GRUs

Update Gate: Determines how much of the past information needs to be passed along to the future.

Reset Gate: Decides how much of the past information to forget.

Simplified Architecture: Compared to Long Short-Term Memory (LSTM) networks, GRUs have a more streamlined structure, often leading to faster training times while maintaining performance.

Source: <https://aws.amazon.com/what-is/recurrent-neural-network/>

```
[2]: train_df.head()
```

```
[2]:      datetime  season  holiday  workingday  weather  temp  atemp  \
0  2011-01-01 00:00:00      1       0          0        1   9.84  14.395
1  2011-01-01 01:00:00      1       0          0        1   9.02  13.635
2  2011-01-01 02:00:00      1       0          0        1   9.02  13.635
3  2011-01-01 03:00:00      1       0          0        1   9.84  14.395
4  2011-01-01 04:00:00      1       0          0        1   9.84  14.395

      humidity  windspeed  casual  registered  count
```

0	81	0.0	3	13	16
1	80	0.0	8	32	40
2	80	0.0	5	27	32
3	75	0.0	3	10	13
4	75	0.0	0	1	1

```
[3]: train_df.dtypes
```

```
[3]: datetime      object
season          int64
holiday         int64
workingday      int64
weather         int64
temp           float64
atemp          float64
humidity        int64
windspeed      float64
casual          int64
registered      int64
count          int64
dtype: object
```

```
[4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   datetime        10886 non-null object
1   season          10886 non-null int64
2   holiday         10886 non-null int64
3   workingday      10886 non-null int64
4   weather         10886 non-null int64
5   temp           10886 non-null float64
6   atemp          10886 non-null float64
7   humidity        10886 non-null int64
8   windspeed      10886 non-null float64
9   casual          10886 non-null int64
10  registered      10886 non-null int64
11  count           10886 non-null int64
dtypes: float64(3), int64(8), object(1)
memory usage: 1020.7+ KB
```

```
[5]: test_df_raw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6493 entries, 0 to 6492
```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	datetime	6493 non-null	object
1	season	6493 non-null	int64
2	holiday	6493 non-null	int64
3	workingday	6493 non-null	int64
4	weather	6493 non-null	int64
5	temp	6493 non-null	float64
6	atemp	6493 non-null	float64
7	humidity	6493 non-null	int64
8	windspeed	6493 non-null	float64

dtypes: float64(3), int64(5), object(1)

memory usage: 456.7+ KB

```
[6]: # === 2. Feature Engineering: Extract datetime info ===
def add_time_features(df):
    df['datetime'] = pd.to_datetime(df['datetime']) # Convert datetime string
    # to timestamp
    df['hour'] = df['datetime'].dt.hour # Hour of the day (0-23)
    df['day'] = df['datetime'].dt.dayofweek # Day of the week
    # (0=Monday)
    df['month'] = df['datetime'].dt.month # Month (1-12)
    df['year'] = df['datetime'].dt.year # Year (2011 or 2012)
    return df

# Apply time feature extraction
train_df = add_time_features(train_df)
test_df = add_time_features(test_df_raw)
```

```
[7]: train_df.describe()
```

	datetime	season	holiday	\
count	10886	10886.000000	10886.000000	
mean	2011-12-27 05:56:22.399411968	2.506614	0.028569	
min	2011-01-01 00:00:00	1.000000	0.000000	
25%	2011-07-02 07:15:00	2.000000	0.000000	
50%	2012-01-01 20:30:00	3.000000	0.000000	
75%	2012-07-01 12:45:00	4.000000	0.000000	
max	2012-12-19 23:00:00	4.000000	1.000000	
std	NaN	1.116174	0.166599	

	workingday	weather	temp	atemp	humidity	\
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000	
mean	0.680875	1.418427	20.23086	23.655084	61.886460	
min	0.000000	1.000000	0.82000	0.760000	0.000000	
25%	0.000000	1.000000	13.94000	16.665000	47.000000	

50%	1.000000	1.000000	20.50000	24.240000	62.000000
75%	1.000000	2.000000	26.24000	31.060000	77.000000
max	1.000000	4.000000	41.00000	45.455000	100.000000
std	0.466159	0.633839	7.79159	8.474601	19.245033

	windspeed	casual	registered	count	hour \
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000
mean	12.799395	36.021955	155.552177	191.574132	11.541613
min	0.000000	0.000000	0.000000	1.000000	0.000000
25%	7.001500	4.000000	36.000000	42.000000	6.000000
50%	12.998000	17.000000	118.000000	145.000000	12.000000
75%	16.997900	49.000000	222.000000	284.000000	18.000000
max	56.996900	367.000000	886.000000	977.000000	23.000000
std	8.164537	49.960477	151.039033	181.144454	6.915838

	day	month	year
count	10886.000000	10886.000000	10886.000000
mean	3.013963	6.521495	2011.501929
min	0.000000	1.000000	2011.000000
25%	1.000000	4.000000	2011.000000
50%	3.000000	7.000000	2012.000000
75%	5.000000	10.000000	2012.000000
max	6.000000	12.000000	2012.000000
std	2.004585	3.444373	0.500019

```
[8]: print(train_df.dtypes)
```

```
datetime      datetime64[ns]
season         int64
holiday        int64
workingday     int64
weather        int64
temp           float64
atemp          float64
humidity       int64
windspeed      float64
casual         int64
registered     int64
count          int64
hour           int32
day            int32
month          int32
year           int32
dtype: object
```

```
[9]: print(test_df.dtypes)
```

```
datetime      datetime64[ns]
```

```

season                int64
holiday               int64
workingday            int64
weather               int64
temp                  float64
atemp                  float64
humidity              int64
windspeed              float64
hour                  int32
day                   int32
month                 int32
year                  int32
dtype: object

```

```

[10]: # === 3. Define input features and target ===
features = ['hour', 'day', 'month', 'year', 'temp', 'humidity', 'windspeed']
target = 'count' # This is what we want to predict

#===3.5 Dropping unnecessary features
dropFeatures = ["datetime"]
train_df = train_df.drop('casual', axis=1)
train_df = train_df.drop('registered', axis=1)
train_df = train_df.drop(dropFeatures,axis=1)
test_df = test_df.drop(dropFeatures,axis=1)

# === 4. Normalize features using MinMaxScaler ===
scaler = MinMaxScaler()
X_train_all = scaler.fit_transform(train_df[features]) # Fit only on train data
X_test_all = scaler.transform(test_df[features]) # Transform test data
↳using same scaler

# Save target variable separately
y_train_all = np.log1p(train_df['count'].values)
#The following is for ensemble and tree methods

```

```

[11]: # === 6. Create time-series sequences for GRU ===
def create_sequences(X, y=None, look_back=24):
    Xs, ys = [], []
    for i in range(len(X) - look_back):
        Xs.append(X[i:i + look_back])
        if y is not None:
            ys.append(y[i + look_back])
    return np.array(Xs), np.array(ys) if y is not None else np.array(Xs)

look_back = 6 # Using past 6 time value

```

```
X_train_seq, y_train_seq = create_sequences(X_train_all, y_train_all, look_back)

#y_train_seq = np.log1p(y_train_seq)
```

```
[12]: # === 7. Split training data into train/validation ===
X_train_final, X_val, y_train_final, y_val = train_test_split(
    X_train_seq, y_train_seq, test_size=0.2, random_state=42, shuffle=False
)
# Note: shuffle=False preserves time order, which is critical for time series

# === 8. Build deeper GRU model ===
model = Sequential()
model.add(Bidirectional(GRU(512, return_sequences=True),
    ↪input_shape=(look_back, len(features))))
#model.add(Dropout(0.05))
model.add(Bidirectional(GRU(128)))
model.add(Dense(1)) # Output: predicted log(count)
model.compile(optimizer='adam', loss='mean_squared_error')

# === 10. Add early stopping ===
early_stop = EarlyStopping(monitor='val_loss', patience=10,
    ↪restore_best_weights=True)

# === 11. Train model with validation data ===
history = model.fit(
    X_train_final, y_train_final,
    validation_data=(X_val, y_val),
    epochs=1000,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/rnn/bidirectional.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)
```

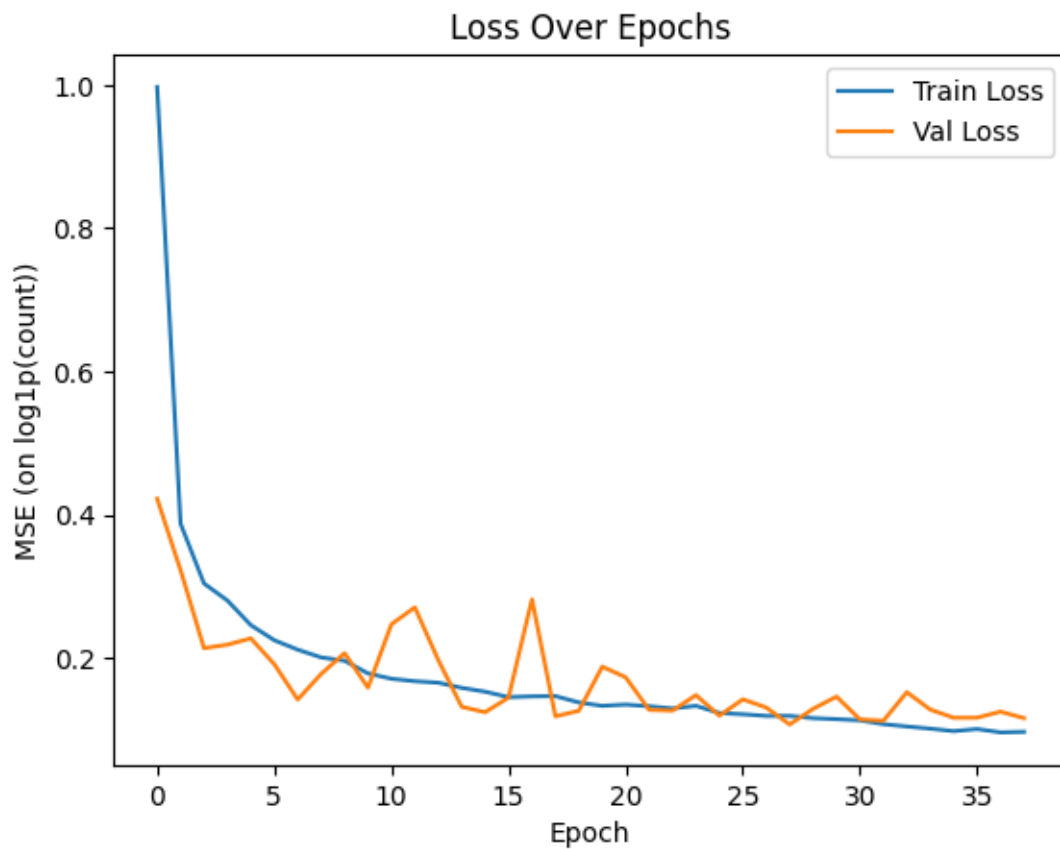
```
Epoch 1/1000
272/272          7s 12ms/step -
loss: 2.1393 - val_loss: 0.4221
Epoch 2/1000
272/272          2s 8ms/step -
loss: 0.3986 - val_loss: 0.3224
Epoch 3/1000
272/272          2s 8ms/step -
```

loss: 0.3169 - val_loss: 0.2133
Epoch 4/1000
272/272 2s 8ms/step -
loss: 0.2870 - val_loss: 0.2182
Epoch 5/1000
272/272 2s 8ms/step -
loss: 0.2639 - val_loss: 0.2270
Epoch 6/1000
272/272 2s 8ms/step -
loss: 0.2214 - val_loss: 0.1910
Epoch 7/1000
272/272 2s 8ms/step -
loss: 0.2117 - val_loss: 0.1415
Epoch 8/1000
272/272 2s 8ms/step -
loss: 0.2023 - val_loss: 0.1773
Epoch 9/1000
272/272 2s 8ms/step -
loss: 0.1893 - val_loss: 0.2063
Epoch 10/1000
272/272 2s 8ms/step -
loss: 0.1852 - val_loss: 0.1581
Epoch 11/1000
272/272 2s 8ms/step -
loss: 0.1641 - val_loss: 0.2467
Epoch 12/1000
272/272 2s 8ms/step -
loss: 0.1696 - val_loss: 0.2703
Epoch 13/1000
272/272 2s 8ms/step -
loss: 0.1669 - val_loss: 0.1968
Epoch 14/1000
272/272 2s 8ms/step -
loss: 0.1586 - val_loss: 0.1314
Epoch 15/1000
272/272 2s 8ms/step -
loss: 0.1494 - val_loss: 0.1238
Epoch 16/1000
272/272 2s 8ms/step -
loss: 0.1448 - val_loss: 0.1443
Epoch 17/1000
272/272 2s 8ms/step -
loss: 0.1433 - val_loss: 0.2815
Epoch 18/1000
272/272 2s 8ms/step -
loss: 0.1486 - val_loss: 0.1181
Epoch 19/1000
272/272 2s 8ms/step -

loss: 0.1351 - val_loss: 0.1258
Epoch 20/1000
272/272 2s 8ms/step -
loss: 0.1289 - val_loss: 0.1875
Epoch 21/1000
272/272 2s 8ms/step -
loss: 0.1409 - val_loss: 0.1726
Epoch 22/1000
272/272 2s 8ms/step -
loss: 0.1280 - val_loss: 0.1272
Epoch 23/1000
272/272 2s 8ms/step -
loss: 0.1249 - val_loss: 0.1261
Epoch 24/1000
272/272 2s 8ms/step -
loss: 0.1352 - val_loss: 0.1477
Epoch 25/1000
272/272 2s 8ms/step -
loss: 0.1274 - val_loss: 0.1189
Epoch 26/1000
272/272 2s 8ms/step -
loss: 0.1173 - val_loss: 0.1418
Epoch 27/1000
272/272 2s 8ms/step -
loss: 0.1155 - val_loss: 0.1304
Epoch 28/1000
272/272 2s 8ms/step -
loss: 0.1228 - val_loss: 0.1066
Epoch 29/1000
272/272 2s 8ms/step -
loss: 0.1151 - val_loss: 0.1285
Epoch 30/1000
272/272 2s 8ms/step -
loss: 0.1164 - val_loss: 0.1456
Epoch 31/1000
272/272 2s 8ms/step -
loss: 0.1089 - val_loss: 0.1137
Epoch 32/1000
272/272 2s 8ms/step -
loss: 0.1046 - val_loss: 0.1119
Epoch 33/1000
272/272 2s 8ms/step -
loss: 0.1023 - val_loss: 0.1518
Epoch 34/1000
272/272 2s 8ms/step -
loss: 0.1030 - val_loss: 0.1275
Epoch 35/1000
272/272 2s 8ms/step -


```
loss: 0.0944 - val_loss: 0.1161
Epoch 36/1000
272/272          2s 8ms/step -
loss: 0.0918 - val_loss: 0.1162
Epoch 37/1000
272/272          2s 8ms/step -
loss: 0.0925 - val_loss: 0.1246
Epoch 38/1000
272/272          2s 8ms/step -
loss: 0.0930 - val_loss: 0.1156
```

```
[13]: # === 12. Plot training & validation loss ===
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('MSE (on log1p(count))')
plt.legend()
plt.show()
```



```
[14]: #We can also check root mean square error if need be

def rmsle(y, y_,convertExp=True):
    log1 = np.nan_to_num(np.array([np.log(v + 1) for v in y]))
    log2 = np.nan_to_num(np.array([np.log(v + 1) for v in y_]))
    calc = (log1 - log2) ** 2
    return np.sqrt(np.mean(calc))
```

```
[15]: # === 13. Prepare test sequences using last part of train ===
X_test_seq = []
seed_seq = X_train_all[-look_back:].tolist()
for row in X_test_all:
    seed_seq.append(row)
    seq = seed_seq[-look_back:]
    X_test_seq.append(seq)
X_test_seq = np.array(X_test_seq)

# === 14. Predict and inverse log1p ==
test_preds = model.predict(X_test_seq).flatten()
```

203/203 1s 3ms/step

```
[16]: # === 13. Submission file ===
datetimecol = test_df_raw["datetime"]
submission = pd.DataFrame({
    "datetime": datetimecol,
    "count": [max(0, x) for x in np.exp(test_preds)]
})
submission.to_csv('bike_predictions3.csv', index=False)
```

```
[18]: from sklearn.ensemble import GradientBoostingRegressor
gbm = GradientBoostingRegressor(n_estimators=4000)
gbm.fit(X_train_all, y_train_all)
preds = gbm.predict(X=X_train_all)
print ("RMSLE Value For Gradient Boost: ", rmsle(np.exp(y_train_all), np.
    ↪exp(preds), False))
```

RMSLE Value For Gradient Boost: 0.2256738796033782

```
[19]: test_preds2 = gbm.predict(X=X_test_all)
```

```
[20]: submission2 = pd.read_csv('/kaggle/input/bike-sharing-demand/sampleSubmission.
    ↪csv')
datetimecol = test_df_raw["datetime"]
submission2 = pd.DataFrame({
    "datetime": datetimecol,
    "count": [max(0, x) for x in np.exp(test_preds2)]
})
```

```
submission2.to_csv('gbr_estt.csv', index=False)
```

```
[21]: from sklearn.ensemble import RandomForestRegressor
rfModel = RandomForestRegressor(n_estimators=100)
rfModel.fit(X_train_all, y_train_all)
preds = rfModel.predict(X=X_train_all)
print ("RMSLE Value For Random Forest: ",rmsle(np.exp(y_train_all), np.
↪exp(preds), False))
```

RMSLE Value For Random Forest: 0.11502367206001117

```
[22]: test_pred = rfModel.predict(X=X_test_all)
```

```
[23]: submission3 = pd.read_csv('/kaggle/input/bike-sharing-demand/sampleSubmission.
↪csv')
datetimecol = test_df_raw["datetime"]
submission3 = pd.DataFrame({
    "datetime": datetimecol,
    "count": [max(0, x) for x in np.exp(test_pred)]
})
submission3.to_csv('bikeseparate_rf.csv', index=False)
```

Random forest gave the best result as a public score 0.39. GRU wasn't the best. However I haven't utilized grid search for the best hyperparameters

```
[ ]:
```